

Greedy Algorithms and Huffman Coding

Henry Z. Lo

June 10, 2014

1 Greedy Algorithms

1.1 Change making problem

Problem 1. *You have quarters, dimes, nickels, and pennies. Given some amount, n , provide the least number of coins which sum up to n .*

For example, let $n = 15$. There are four different coin combinations to get 15¢ (see Figure 1). The optimal combination is one dime and one nickel.

It turns out for this problem, we can just pick the largest coin value possible, and repeat until done. Because this algorithm takes the largest possible step to its goal every time (e.g. from 0 to 15), it is called *greedy*. This produces the optimal solution for this particular set of coins, but does not work all the time.

Problem 2. *You have 1¢, 7¢, and 9¢ coins. Given some amount, n , provide the least number of coins which sum up to n .*

If we let $n = 15$ again, then the first greedy choice takes the 9¢ coin. There are 6¢ remaining, which must be made up for with 1¢ coins. This yields a 7 coin solution. The optimal configuration is 7¢, 7¢, 1¢.



Figure 1: The four ways of getting 15 cents. The dime and nickel combination is optimal.

1.2 Greedy methods

Greedy methods apply to combinatorial optimization problems. In these problems, we try to find some optimal set or combination of elements to minimize or maximize something given constraints. In the previous case, the elements are coins, the constraints are that the coins sum to n , and we want to minimize the number of coins.

These problems pop up all over computer science, so we will spend some time studying them.

1.2.1 Problems solved by greedy algorithms

To determine if a greedy algorithm works for a problem, we need to show that the problem has the following properties:

- *Greedy choice property*: an optimal solution contains the greedy choice.
- *Optimal substructure*: optimal solutions consists of an optimal subsolution, plus an optimal choice.

Change making contains optimal substructure. Consider an optimal solution for n . Then the coins making up $n - c$, where c is an optimal choice, must have been optimal. If it wasn't, then you can find some set to get less coins to make up $n - c$, and the original optimal solution was not optimal.

Change making does not contain the greedy choice property in general - as we've seen, sometimes the optimal solution does not contain the greedy choice. However, in the US currency, it does. If $n \geq 25$, the optimal solution contains a quarter; if $n \geq 10$, it contains a dime, etc.

We will not emphasize mathematical proofs in this course, but know that there is a method for determining when greedy algorithms can be applied.

1.2.2 Properties of greedy algorithms

Despite their limitations, it's worth noting greedy algorithms because when they do solve a problem, they are extremely efficient. There is no backtracking, and they are easy to program. Consider the following programs for the change-making problem.

Greedy solution:

```
sum, numcoins = 0
coins = [25,10,5,1]
while (sum < n)
    numcoins++
    for c in coins:
        if (n-sum >= c)
            sum = sum+c
```

There is no algorithm more efficient. You would have to build a tree-like structure for exhaustively searching the problem space.




Figure 2: A set of activities with finish and end times. If we want to perform the most activities, A1, A4, A8, and A11 is optimal.

1.2.3 Drawbacks

Greedy methods always take the largest step towards its goals, but in doing so may miss some better combinations, as seen in the second problem. They do not take into account the future (if I take 7¢ instead of 9¢ now, then I can take another big coin in my next step).

However, even when greedy algorithms don't give the best result, sometimes the greedy choice is a good *heuristic* for getting reasonable solutions (it can also be a very bad heuristic). Some problems are intractable, and greedy algorithms can give decent answers in a reasonable time.

1.3 Activity selection

Problem 3. *There are n activities, each with start times s_i and finish times f_i . How do you pick a set of non-overlapping activities so that you perform the most possible activities?*

The solution, surprisingly, consists of picking the earliest ending activities, removing the overlapping activities, and repeat (see Figure 2).




Figure 4: Huffman algorithm for combining characters a1, a2, a3, and a4. Red numbers represent frequency, and yellow numbers represent code.

2.2 Huffman codes

Huffman codes solve the problem of finding the optimal encoding. Note that the code depends on the text being compressed.

We want to have shorter codes / shorter paths / less parents for characters with very high counts, but we are alright with having many parents for characters with less counts.

1. Create a leaf node for each character, put them in priority queue in order of counts.
2. While the queue has more than one character:
 - Remove two nodes with lowest counts.
 - Merge them to have one parent, with a cumulative sum.
 - Add this node to the queue.
3. Last node will be the root node.

See Figure 4 for a visual.

Since efficient priority queue data structures require $O(\log n)$ time per insertion, and a tree with n leaves has $2n - 1$ nodes, this algorithm operates in $O(n \log n)$ time, where n is the number of characters.

2.3 Greediness of Huffman codes

Huffman codes are greedy in the sense that the least frequent codes are given the longest paths in the tree.

Greedy choice property: if x and y have the two lowest counts, then the optimal algorithm will have them at the same path length, i.e. their codes are two of the longest.

Optimal substructure: every subtree of a Huffman code defines an optimal code for its leaf characters.